

Homework: Analysis of Media and Semantic Forensics in Scientific Literature

Due: Friday, March 6, 2020 12pm PT

1. Overview

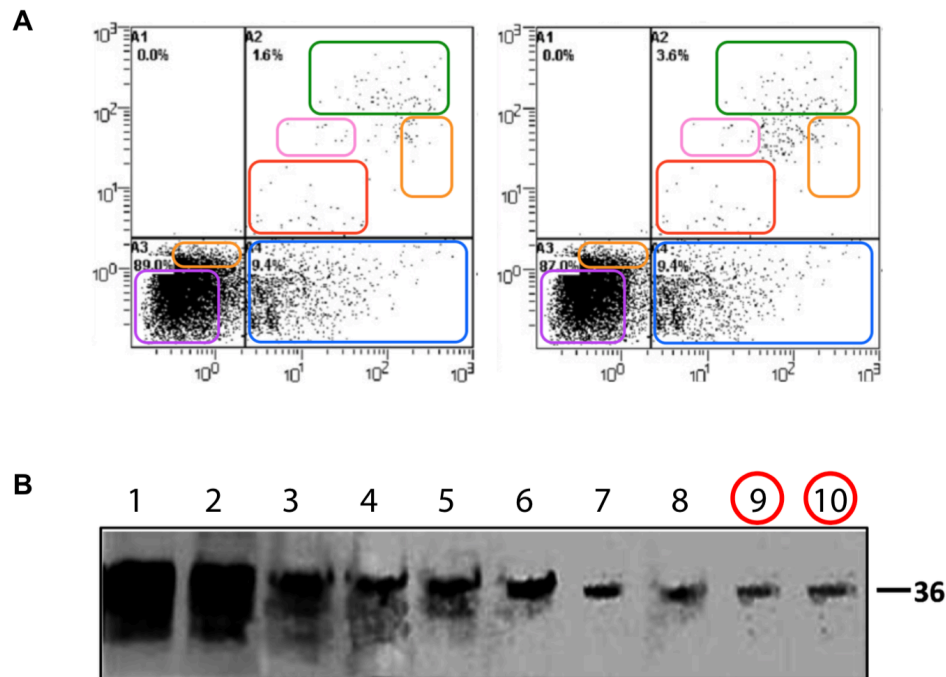


Figure 1: Potentially problematic Media manipulations present in biomedical research papers. A and B show areas of duplicity in included media in papers.

<https://mbio.asm.org/content/7/3/e00809-16.abstract>

In this assignment we will explore several of the topics discussed in the early portion of class – Big Data – MIME types and their taxonomy – Data Similarity – and so forth. To do this, we will leverage the dataset highlighted in Figure 1 – a set of 200 papers (out of 800) from biomedical and scientific literature identified in the Bik et al paper referenced above that have an "endpoint" conclusion. Endpoint papers are papers from the Bik et al 2017 dataset that resulted in either a retraction, a correction, or a no-action (a statement by the journal that they will not investigate because the paper is either too old, or they do not see any problem). Some examples of media manipulation in papers are summarized graphically shown in Figure 1.

Media manipulation is an increasing concern in scientific and open literature since researchers like Bik et al have seen a tremendous uptick in papers with media manipulations and potential problems over the last decade. Figure 2 plots this growth:





Figure 2: Growth of papers with problematic or manipulated images 1995-2014

Additionally, one of the challenges too if you look deeper at the data is that there has been a geographic uptick in the author institutions and affiliations associated with particular countries that seem to be responsible for the uptick when you plot author



affiliation / country against fraction of published papers and fraction of media manipulations. Figure 3 has the associated data and presents it visually.

As can be gleaned from the data, the manipulations occur most by authors from The United States and China and then there is a set of countries with similar properties with respect to number of high impact papers and possible manipulations. Bik et al have captured a dataset that records information about papers including Authors, Paper Title, Citation, Digital Object Identifier (DOI), Year, Month, Classification Label (0-3) referring to three major categories: simple duplications, duplications with repositioning, and duplications with alteration and also two potentially problematic areas that may not be manipulations but cause issues (Cuts & Beautification). The dataset also includes a text description of what was found (Findings), if the paper was reported to the journal or not, and whether it was retracted, or a correction issued or not, and finally whether no action was taken and what date (if any) any action was completed on. The data is formatted according to the referenced schema which will be provided to you and is in TSV (TSV) format (MIME type: text/tab-separated-values).

As you can imagine, there are plenty of things we could do with this data. Some of them fall into to the realm of the paper we discussed from IEEE Computer in 2013 regarding “Big Data’s Unintended Consequences”. If you recall, one of the key points of this paper was that companies and governments were moving away from data silos, and instead focused on how data could be *joined* together to form even more comprehensive, statistically relevant, and accurate data on *everyone* – in short, increasing its value, and potentially its veracity, in addition to its volume, and potentially its variety.

Another topic frequently discussed in class thus far has been the framework of the “Five V’s” – volume, velocity, variety, veracity and value. We have done several class/group activities so far and thought about many datasets in the real world and how to classify and leverage them along this framework. From judicial data about inmates and prisoners, to health data, to Twitter data, to restaurant review data, there have been extremely useful discussions and points made. You will leverage those discussions, class lectures, and material discussed, in this assignment.

2. Objective

Looking at the Bik et al media manipulation data, you may ask yourselves: “what other data is available that could be joined with this information” to affect its Five V’s – intentionally, or unintentionally. For example, consider if you could automatically pull down a lot more information about a particular author of a paper and her co-authors – their rate of publication, number of students in their lab, information about what other journals they have published in, etc. This is an interesting question, and will form the first part of your assignment which will involve finding additional information about each author for each of the provided publications and collecting and joining it to the Bik dataset. You will add new features to the dataset of “Lab Size (number of students)”, “Publication Rate”, “Other Journals Published In” and some information about “First Author” including “Affiliation University”, “Duration of Career (Years)”, highest degree

obtained (e.g., “PhD”, “MS”) and “Degree Area” (e.g., Computer Science). Perhaps there is a pattern that will emerge, for example you could posit that those with a Masters in Computer Science, with 50 years of experience and 100 students in the lab, may not be critically reviewing papers published in biomedical journals. Or not. These will be the interesting research questions that you will investigate!

What other datasets could you join the Bik et al data to? For example, what about census population per closest city to the first author’s institutional affiliation? What about restaurant data indicating open hours and close times for nearby coffee shops - maybe the students in the lab are up late at night and not paying attention when they paste in graphics in the papers? Finally, what about animal population data in the closest county, city, metropolitan area to the author’s lab, to give you an idea of the mice population or rabbit or experimental animal population – perhaps there is duplication because all of the animals are cousins and their cellular data is in fact the same. You could also look at associated conference or journal submission times for the journals or conferences in question, see if the submission times are all late at night....etc.

You will choose at least three publicly accessible datasets along these lines to join the Bik et al data to, and you must add at least three new features per dataset that you join. The datasets you select may not all belong to the same MIME top level type – that is – you must pick a different MIME top level type for each of the three datasets you are joining to this Bik et al problematic papers dataset.

Once the data is joined properly, you will explore the combined dataset using Apache Tika and an associated Python library called Tika-Similarity. Using Tika Similarity, you can evaluate data *similarity* (as discussed during the Deduplication lecture in class; and also during data forensics discussions). Tika similarity will allow you to explore and test different distance metrics (Edit-Distance; Jaccard similarity; Cosine similarity, etc.). And it will give you an idea of how to cluster data, and finally it will let you visualize the differences between different clusters in your new combined dataset. So you can figure out how similar papers with problem areas are within the data, and ask questions of your new augmented Bik et al dataset. For example, you may ask, how many papers with similar media manipulations came from first authors with more than 50 students in the lab, with small mice populations who where all cousins, and in which the students in the lab have access to coffee shops that keep them up all night and thus they are tired when clicking submit on the papers?

The assignment specific tasks will be specified in the following section.

3. Tasks

1. Download and install Apache Tika
 - a. Chapter 2 in your book covers some of the basics of building the code, and additionally, see <http://tika.apache.org/1.23/gettingstarted.html>
 - b. Install Tika-Python, you can pip install tika to get started.
 - i. Read up on Tika Python here:
<http://github.com/chrismattmann/tika-python>

2. Download and install D3.js
 - a. Visit <http://d3js.org/>
 - b. Review Mike Bostock's Visual Gallery Wiki
 - i. <https://github.com/mbostock/d3/wiki/Tutorials>
3. Download the Bik et al dataset from the Dropbox
 - a. <https://www.dropbox.com/s/11qrb6f05myd9a4/Bik%20dataset%20-%20papers%20with%20endpoint%20reached.tsv?dl=0>
 - b. Make a copy of the original dataset (because you are going to modify/add to it in this assignment)
4. Begin by adding additional features to the data about the author's lab size, affiliation, etc:
 - a. Lab Size (number of students)
 - b. Publication Rate
 - c. Other Journals Published In
 - d. Information about First Author including
 - i. Affiliation University
 - ii. Duration of Career (Years)
 - iii. Highest degree obtained (e.g., "PhD", "MS")
 - iv. Degree Area" (e.g., Computer Science)
 - e. Write a Python program to collect the above information for each of the 200 papers in the dataset:
 - i. **EXTRA CREDIT:** Use USC Data Science's Sparkler crawler to collect this data
 1. <http://github.com/USCDataScience/sparkler>
5. Identify at least three other datasets, each of different top level MIME type (can't all be e.g., text/*)
 - a. Check out places including: <https://catalog.data.gov/dataset> (Data.gov)
 - b. For each dataset, develop a Python program to join the data to your Bik et al papers dataset
 - i. For each non text/* dataset, be prepared to describe how you featurized the dataset
 - c. Each dataset that you join must contribute at least three features (in addition to the features you are adding described in part 4)
 - d. For each feature you add, be prepared to discuss what types of queries it will allow you to answer and also how you computed the feature
6. Download and install Tika-Similarity
 - a. Read the documentation
 - b. You can find Tika Similarity here (<http://github.com/chris mattmann/tika-similarity>)
 - c. Compare Jaccard similarity, edit-distance, and cosine similarity
 - i. Compare and contrast clusters from Jaccard, Cosine Distance, and Edit Similarity – do you see any differences? Why?
 - d. How to the resultant clusters generated highlight the features you extracted? Be prepared to identify this in your report
7. **(EXTRA CREDIT)** Add some new D3.js visualizations to Tika Similarity

- a. Currently Tika Similarity only supports Dendrogram, Circle Packing, and combinations of those to view clusters, and relative similarities between datasets
- b. Consider adding
 - i. Feature related visualizations, e.g., time series, bar charts, plots
 - ii. Add functionality in a generic way that is not specific to your dataset
 - iii. See gallery here: <https://github.com/d3/d3/wiki/Gallery>
 - iv. Contributions will be reviewed as Pull Requests in a first come, first serve basis (check existing PRs and make sure you aren't duplicating what some other group has done)

4. Assignment Setup

4.1 Group Formation

You can work on this assignment in groups sized at minimum 2, and maximum 5. You may reuse your existing groups from discussion in class. Please email your group ONE time (from the main point of contact) to Yash Shah, our TA after class on Thursday, February 13. If you have questions contact Yash via his e-mail address on the class website with the subject: INF 550: Team Details.

4.2 Bik et al papers dataset

Access to the data is provided in the Dropbox. The dataset itself is 96k. You may want to distribute the data between your team-mates since the data is fairly small (for now).

4.3 Downloading and Installing Apache Tika

The quickest and best way to get Apache Tika up and running on your machine is to grab the tika-app.jar from: <http://tika.apache.org/download.html>. You should obtain a jar file called tika-app-1.23.jar. This jar contains all of the necessary dependencies to get up and running with Tika by calling it your Java program.

Documentation is available on the Apache Tika webpage at <http://tika.apache.org/>. API documentation can be found at <http://tika.apache.org/1.23/api>.

Since you will be using Tika Python, you will want to read up on the Tika REST API, here: <https://wiki.apache.org/tika/TikaJAXRS>. The Tika Python library is a robust REST client to the Java-side REST API.

You can also get more information about Tika by checking out the book written by Professor Mattmann called "Tika in Action", available from: <http://manning.com/mattmann/>.

5. Report

Write a short 4 page report describing your observations, i.e. what you noticed about the dataset as you completed the tasks. What questions did your new joined datasets allow you to answer about the Bik et al papers previously unanswered? What clusters were revealed? What similarity metrics produced more (in your opinion) accurate measurements? Why? What did the additional datasets suggest about “unintended consequences” related to media forensics data? You should also clearly explain which datasets you used to join the problematic papers data and how you extracted the new features from each dataset.

Thinking more broadly, do you have enough information to answer the following:

1. Are there clusters of authors with similar media manipulations that exhibit at least 3 similarities between the students and properties of their lab?
2. Does staying up late at night matter?
3. Does the animal population for available specimens in the area influence the type of manipulations?
4. What do animal population demographics tell us about the institutions in which media manipulations occur?
 - a. Densely populated? Sparsely populated?
5. What insights do the “indirect” features you extracted tell us about the data?
6. What clusters of papers made the most sense? Why?

Also include your thoughts about Apache Tika – what was easy about using it? What wasn't?

6. Submission Guidelines

This assignment is to be submitted **electronically, by 12pm PT** on the specified due date, via Gmail inf550spring2020@gmail.com. Use the subject line: INF 550: Mattmann: Spring 2020: BIGDATA Homework: Team XX. So if your team was team 15, you would submit an email to inf550spring2020@gmail.com with the subject “INF 550: Mattmann: Spring 2020: BIGDATA Homework: Team 15” (no quotes). **Please note only one submission per team.**

- All source code is expected to be commented, to compile, and to run. You should have at least a few Python scripts that you used to join three other datasets, and what you used to extract additional features.
- Include your updated dataset TSV. We will provide a Dropbox location for you to upload to.
- Also prepare a readme.txt containing any notes you'd like to submit.
- If you used external libraries other than Tika Python and Tika Similarity, you should include those jar files in your submission, and include in your readme.txt a detailed explanation of how to use these libraries when compiling and executing your program.
- Save your report as a PDF file (TEAM_XX_BIGDATA.pdf) and include it in your submission.
- Compress all of the above into a single zip archive and name it according to the following filename convention:
TEAM_XX_INF550_HW_BIGDATA.zip
Use only standard zip format. Do **not** use other formats such as zipx, rar, ace, etc.
- If your homework submission exceeds the Gmail's 25MB limit, upload the zip file to Google drive and share it with inf550spring2020@gmail.com.

Important Note:

- Make sure that you have attached the file the when submitting. Failure to do so will be treated as non-submission.
- Successful submission will be indicated in the assignment's submission history. We advise that you check to verify the timestamp, download and double check your zip file for good measure.
- Again, please note, only **one submission per team**. Designate someone to submit.

6.1 Late Assignment Policy

- -10% if submitted within the first 24 hours
- -15% for each additional 24 hours or part thereof